

GB300 NVL72 Cluster.

576 GB300 Blackwell Ultra GPUs available February 2027. 165.9 TB HBM3e. One dedicated cluster: no shared tenancy, no throttling. Located in Canada. Optional: ramp up to a 2,304-GPU cluster by August 2027.

576 GPUs GB300	165.9 TB HBM3E MEMORY	8 SUs NVL72 RACKS	1.8 TB/s NVLINK 5 PER GPU	288 CPUs INTEGRATED GRACE
--------------------------	---------------------------------	-----------------------------	-------------------------------------	-------------------------------------

CLUSTER ARCHITECTURE

Total GPUs	576 × NVIDIA Blackwell Ultra
Scalable Units	8 × GB300 NVL72 racks
GPUs per SU	72, single NVLink 5 domain
Compute Trays	144 total (18 per SU)
Grace CPUs	288 integrated (2 per tray)
NVLink Switch Trays	72 total (9 per SU)
Target Ready	February 25, 2027

MEMORY

HBM3e per GPU	≥ 288 GB
HBM3e per SU	≈ 20.7 TB
HBM3e total	≈ 165.9 TB
Grace Host Mem	Fully populated OEM baseline
Local Storage	OEM boot + NVMe data cache

NETWORKING & FABRIC

Scale-Up	NVLink 5, 1.8 TB/s bidir. per GPU
Scale-Out	Quantum-X800 InfiniBand, rail-optimized 8-SU
Host NICs	4 × ConnectX-8 per tray (800 Gb/s/GPU)
DPU	BlueField-3-class per tray
Traffic Sep.	GPU east-west fully isolated
OOB Mgmt	Dedicated network, NTP/PTP sync

POWER & FACILITY

Power per Rack	120–142 kW OEM profile
Headroom	100–110% of signed OEM profile
Cooling	Liquid-cooled, OEM-spec chemistry
Throttling	No persistent cap on performance
Facility Valid.	OEM-authorized review at delivery
Maintenance	≥ 10 business days notice

ACCESS & OBSERVABILITY

- Full bare-metal access: PXE image deploy, your OS
- DCGM telemetry: clocks, temps, power, ECC, NVLink
- Fabric metrics: link speed, error & congestion counters
- 90-day telemetry retention, 12-month incident records
- 24/7 incident intake with defined severity SLAs

ACCEPTANCE & QUALITY

- Signed ATM: objective pass/fail before we ship
- 24-hour full-load soak per SU before handover
- Extended DCGM diagnostics across all 576 GPUs
- HBM bandwidth, NVLink topology & NCCL benchmarks
- 10 business-day customer test window per Tranche
- Rack-by-rack: one failed SU won't block others

OPTIONAL · AVAILABLE BY AUGUST 2027

Scale to 2,304 GPUs.

Expand from the initial 576-GPU cluster to a 2,304-GPU fleet across 32 GB300 NVL72 Scalable Units. Same hardware generation, same fabric architecture, same acceptance standards.

Aug 2027

TARGET READY